






2M batch LR=0.04 debug



Readout

1. LR=0.04 is monotonic during training; it is not a run-level divergence.
2. The jagged LR sweep is mainly a final-BPB ordering issue: identity lr=0.04 underperforms both lr=0.02 and lr=0.08 in the same seed-42 sweep.
3. Final validation loss after a fixed token budget is not mathematically convex in LR; optimizer dynamics, warmup/decay schedule, stochasticity, and finite validation noise can create local reversals.
4. Treat single-seed nonconvexity as a signal to add seeds around the minimum, not as a principled optimizer ranking.